This project is based on a project by Jon Shidal designed for the *Operating Systems* course at Washington University in St. Louis.

1   Overview

Like the first project, this project will have 2 parts. In the first part you will get to implement your own user space memory allocator. You will learn the complexities and details of memory allocation and handling free lists first-hand. You will also learn to build a shared dynamic library and you may even think a little bit about the efficiency of your code.

In the second part, you will first become familiar with how xv6 virtualizes memory and then add a couple of additional features to xv6 that many modern operating systems already have.

The projects from here on out will most likely take much more time than the first, so you are welcome to work in teams of up to 2 for this project. Procrastinating on this project (or future projects) will not turn out well for you.

2   Part A: User-space Memory Allocation

For this part of the project, you will be developing your own user space memory allocation library (which will replace the standard memory allocator, e.g. `malloc()` and `free()`).

2.1   Some resources to get started

Read Chapter 14 on the memory API (if you haven't already) to get a grasp on how `malloc` and `free` are called and what they return. Also have a look at their *man* pages.

Read Chapter 17 on free space management for details on how to go about implementing your own allocator. This chapter will help you tremendously on this project.

Note: You will eventually need to add a header to each allocated object much like the one described in Ch. 17. Your header should not exceed 16 bytes in size.

2.2   Details

Memory allocators have a couple of responsibilities. First, the allocator must request memory from the OS for the process's heap. Typically, allocators start with a small heap and ask the OS for more space when needed by the application (this reduces the number of valid page table entries for the given process to only those needed by the process). Allocators may use `sbrk` or `mmap` to ask the OS for additional space. We will simplify our allocator by only asking the OS for memory space once. When we are out of heap space, we are simply out of space.

The second task of an allocator is to manage the heap, generally by maintaining a list of unused chunks of memory. When the process asks to allocate some heap space, the allocator chooses a chunk of free memory and returns it to the process via a pointer to it. When the process no longer needs a memory chunk, the allocator adds the freed memory back to the list of unused memory chunks.

The user space memory allocator is generally included as part of the standard library. It is not part of the OS. It runs fully in the virtual address space of the process, with no concept of pages or physical addresses. Address translation and the mapping of virtual pages to physical page frames is handled entirely by the OS.

Classic `malloc()` and `free()` are defined as follows:

- `void *malloc(size_t size)`: allocates size bytes and returns a pointer to the allocated memory. The memory is not cleared or initialized in any way.

- `void free(void *ptr)`: frees the memory space pointed to by ptr, which must have been returned by a previous call to `malloc()` (or `calloc()` or `realloc()`). Otherwise, or if `free(ptr)` has already been called before, undefined behaviour occurs. If `ptr` is NULL, no operation is performed.

Your implementations of `Mem_Alloc(int size, int policy)` and `Mem_Free(void *ptr)` should follow what `malloc()` and `free()` do; see below for details.

You will also provide a supporting function, `Mem_Dump()`, described below; this routine simply prints which regions are currently free and should be used by you for debugging purposes.

2.3 Program Specifications

For this project, you will be implementing several different routines as part of a shared library. Note that you will not be writing a `main()` routine for the code that you handin (but you should implement one for your own testing). We have provided the prototypes for these functions in the file `mem.h` (which is available in `/home/comp345/proj2a.zip`); you should include this header file in your code to ensure that you are adhering to the specification exactly. You should not change `mem.h` in any way! We now define each of these routines more precisely.

- `int Mem_Init(int size)`: called one time by a process using your routines. `size` is the number of bytes that you should request from the OS using `mmap()`.

- Note that you may need to round up this amount so that you request memory in units of the page size (see the man pages for `getpagesize()`). Note also that you need to use this allocated memory for your own data structures as well; that is, your infrastructure for tracking the mapping from addresses to memory objects has to be placed in this region as well. You are not allowed to `malloc()`, or any other related function, in any of your routines! Similarly, you should not allocate global arrays. However, you may allocate a few global variables (e.g., a pointer to the head of your free list.)

- Return 0 on a success (when call to `mmap` is successful). Otherwise, return -1 and set `m_err` to `ERR_BAD_ARGS`. Cases where `Mem_Init` should return a failure: `Mem_Init` is called more than once; size is less than or equal to 0.

- `void *Mem_Alloc(int size, int policy)`: similar to the library function `malloc()`. `Mem_Alloc` takes as input the size in bytes of the object to be allocated as well as an allocation policy, and returns a pointer to the start of that object. The function returns NULL if there is not enough contiguous free space within the memory allocated by `Mem_Init` to satisfy this request (and sets `m_err` to `ERR_OUT_OF_SPACE`). If `Mem_Init` has not been called yet, you should set `m_err` to `ERR_MEM_UNINITIALIZED` and return NULL.

- You should implement multiple allocation strategies, as described in the book.

  - 0 - Best-Fit
  - 1 - First-Fit
  - 2 - Next-Fit (extra credit - 10 points)

Use the policy number passed to `Mem_Alloc()` to decide which allocation policy to use for the given allocation.

- For performance reasons, `Mem_Alloc()` should return 8-byte aligned chunks of memory. For example, if a user allocates 1 byte of memory, your

2

`Mem_Alloc()` implementation should return 8 bytes of memory so that the next free block will be 8-byte aligned too. To figure out whether you return 8-byte aligned pointers, you could print the pointer this way: `printf("%p", ptr)`. The last digit should be a multiple of 8 (i.e., 0 or 8).

- `int Mem_Free(void *ptr)`: frees the memory object that `ptr` points to. Just like with the standard `free()`, if `ptr` is NULL, then no operation is performed. The function returns 0 on success and -1 otherwise. On failure, set `m_err` to `ERR_INVALID_PTR` prior to returning -1.

- Coalescing: `Mem_Free()` should make sure to coalesce free space. Coalescing rejoins neighboring freed blocks into one bigger free chunk, thus ensuring that big chunks remain free for subsequent calls to `Mem_Alloc()`.

- `void Mem_Dump()`: This is just a debugging routine for your own use. Have it print the regions of free memory to the screen.

You must provide these routines in a shared library named `libmem.so`. Placing the routines in a shared library instead of a simple object file makes it easier for other programmers to link with your code. There are further advantages to shared (dynamic) libraries over static libraries. When you link with a static library, the code for the entire library is merged with your object code to create your executable; if you link to many static libraries, your executable will be enormous. However, when you link to a shared library, the library's code is not merged with your program's object code; instead, a small amount of stub code is inserted into your object code and the stub code finds and invokes the library code when you execute the program. Therefore, shared libraries have two advantages: they lead to smaller executables and they enable users to use the most recent version of the library at run-time. To create a shared library named `libmem.so`, use the following commands (assuming your library code is in a single file `mem.c`):

```
gcc -c -fpic mem.c -Wall -Werror gcc -shared -o libmem.so mem.o
```

To link with this library, you simply specify the base name of the library with `-lmem` and the path so that the linker can find the library with `-L.`:

```
gcc -lmem -L. -o myprogram mymain.c -Wall -Werror
```

Of course, these commands should be placed in a `Makefile`. Before you run `myprogram`, you will need to set the environment variable `LD_LIBRARY_PATH`, so that the system can find your library at run-time. Assuming you always run `myprogram` from this same directory, you can use the command:

```
export LD_LIBRARY_PATH="$LD_LIBRARY_PATH:."
```

2.4 Unix Hints

In this project, you will use `mmap` to map zero'd pages (i.e., allocate new pages) into the address space of the calling process. Note there are a number of different ways that you can call `mmap` to achieve this same goal; we give one example here:

```
// open the /dev/zero device int
fd = open("/dev/zero", O_RDWR);

// size (in bytes) needs to be evenly divisible by the page size
size_t size = ...;
void *ptr = mmap(NULL, size, PROT_READ | PROT_WRITE, MAP_PRIVATE, fd, 0);
if (ptr == MAP_FAILED) { perror("mmap"); exit(1); }

// close the device (don't worry, mapping should be unaffected)
close(fd);
return 0;
```

3   Part B: Explore Memory Virtualization in xv6

In this part of the project, you should get familiar with how xv6 handles memory virtualization. You will have two goals:

1. Create a new system call that takes a virtual address and returns the physical address that it maps to.

2. Check for a program dereferencing a null pointer and throw an exception if so. This is standard in pretty much all operating systems today.

   The order you finish each of these tasks in does not matter. You may find that handling null pointers first may help you understand the memory virtualization structures enough to implement the system call easily.

   Start by reading chapter 1 of the xv6 book here.

3.1   Details

3.1.1   Part 1

You will add a system call named sys_translate to xv6. Translate takes a virtual address as a parameter (a pointer) and returns the physical address that the virtual address maps to. This is similar to the system call you added for project 1, however you will need to look through the code and gain some understanding on xv6 the data structures and techniques xv6 uses to handle address translation.

3.1.2   Part 2

In xv6, the VM system uses a simple two-level page table as discussed in class. As it currently is structured, user code is loaded into the very first part of the address space. Thus, if you dereference a null pointer, you will not see an exception (as you might expect); rather, you will see whatever code is the first bit of code in the program that is running. Try it and see!

   Thus, the first thing you might do is create a program that dereferences a null pointer. It is simple! See if you can do it. Then run it on Linux as well as xv6, to see the difference. Your job here will be to figure out how xv6 sets up a page table. Thus, once again, this project is mostly about understanding the code, and not writing very much. Look at how exec() works to better understand how address spaces get filled with code and in general initialized. That will get you most of the way.

   You should also look at fork() , in particular the part where the address space of the child is created by copying the address space of the parent. What needs to change in there? The rest of your task will be completed by looking through the code to figure out where there are checks or assumptions made about the address space. Think about what happens when you pass a parameter into the kernel, for example; if passing a pointer, the kernel needs to be very careful with it, to ensure you haven't passed it a bad pointer. How does it do this now? Does this code need to change in order to work in your new version of xv6?

   One last hint: you'll have to look at the xv6 makefile as well. In there, user programs are compiled so as to set their entry point (where the first instruction is) to 0. If you change xv6 to make the first page invalid, clearly the entry point will have to be somewhere else (e.g., the next page, or 0x1000). Thus, something in the makefile will need to change to reflect this as well. You should be able to demonstrate what happens when user code tries to access a null pointer. If you do this part correctly, xv6 should trap and kill the process without too much trouble on your part.

4   Submitting

Create a README file for your project. In it, list each member in your group's name and briefly describe:

1. Part 1: what your library does. Compare the allocation strategies you implemented.

2. The changes you made to xv6, including any test/example programs you wrote.

As in project 1, use `git diff` to create a patch with your xv6 changes. Then submit all your files (new files + your patch file) with `handin`.